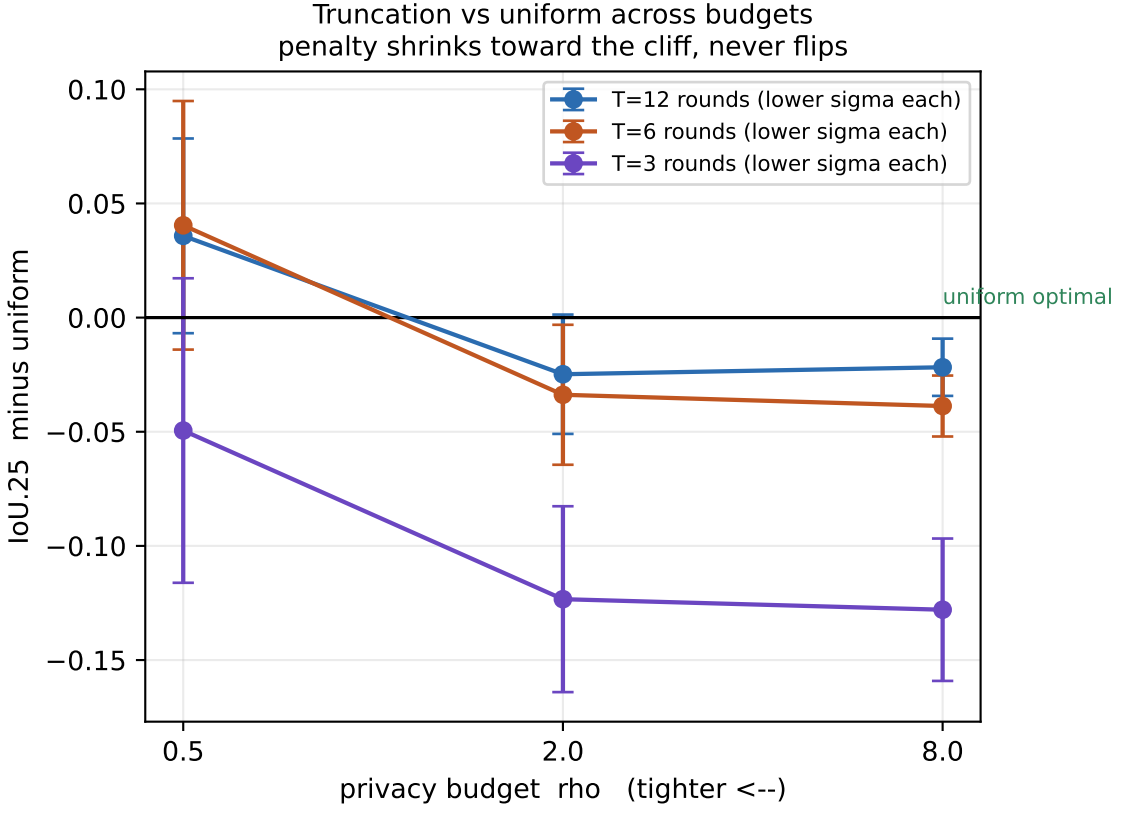
Truncation vs uniform across budgets
penalty shrinks toward the cliff, never flips
T=12 rounds (lower sigma each)
T=6 rounds (lower sigma each)
T=3 rounds (lower sigma each)
uniform optimal
IoU.25 minus uniform
0.10
0.05
0.00
−0.05
−0.10
−0.15
0.5
2.0
8.0
privacy budget rho (tighter <--)

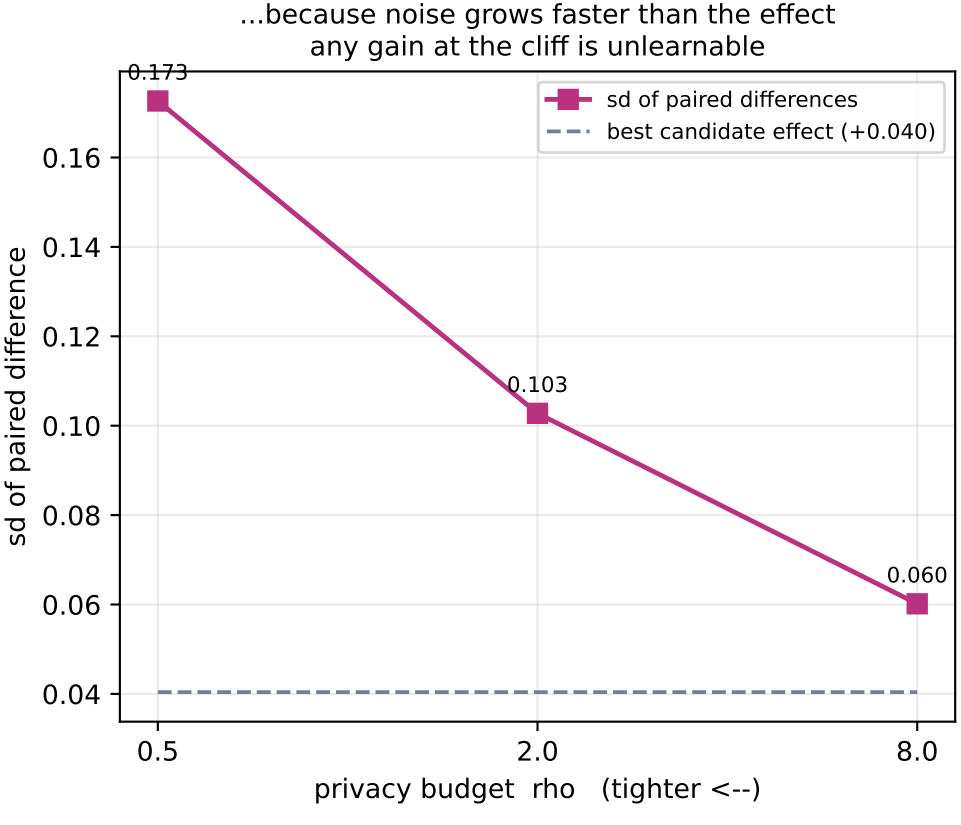
...because noise grows faster than the effect
any gain at the cliff is unlearnable
sd of paired differences
best candidate effect (+0.040)
sd of paired difference
0.173
0.103
0.060
0.16
0.14
0.12
0.10
0.08
0.06
0.04
0.5
2.0
8.0
privacy budget rho (tighter <--)